

# Monitoring

Azure Synapse Analytics provides a rich monitoring experience within the Azure portal to surface insights regarding your data warehouse workload. The Azure portal is the recommended tool when monitoring your data warehouse as it provides configurable retention periods, alerts, recommendations, and customizable charts and dashboards for metrics and logs. The portal also enables you to integrate with other Azure monitoring services such as Azure Monitor (logs) with Log analytics to provide a holistic monitoring experience for not only your data warehouse but also your entire Azure analytics platform for an integrated monitoring experience.

You can monitor active SQL requests using the SQL requests area of the Monitor Hub. This includes details like the pool, submitter, duration, queued duration, workload group assigned, importance, and the request content.

Pipeline runs can be monitored using the Monitor Hub and selecting Pipeline runs. Here you can filter pipeline runs and drill in to view the activity runs associated with the pipeline run and monitor the running of in-progress pipelines.

The execution of Spark applications representing the execution of notebooks and jobs can be monitored within the Monitor Hub, selecting Spark applications. Selecting a Spark application to view its progress and to launch the Spark UI to examine a running Spark job and stage details, or the Spark history server to examine a completed application.

## Exercise 1 - Workload Management

Running mixed workloads can pose resource challenges on busy systems. Solution Architects seek ways to separate classic data warehousing activities (such as loading, transforming, and querying data) to ensure that enough resources exist to hit SLAs.

Synapse SQL pool workload management in Azure Synapse consists of three high-level concepts: Workload Classification, Workload Importance and Workload Isolation. These capabilities give you more control over how your workload utilizes system resources.

Workload importance influences the order in which a request gets access to resources. On a busy system, a request with higher importance has first access to resources. Importance can also ensure ordered access to locks.


Workload isolation reserves resources for a workload group. Resources reserved in a workload group are held exclusively for that workload group to ensure execution. Workload groups also allow you to define the amount of resources that are assigned per request, much like resource classes do. Workload groups give you the ability to reserve or cap the amount of resources a set of requests can consume. Finally, workload groups are a mechanism to apply rules, such as query timeout, to requests.

### Task 1 - Workload Importance

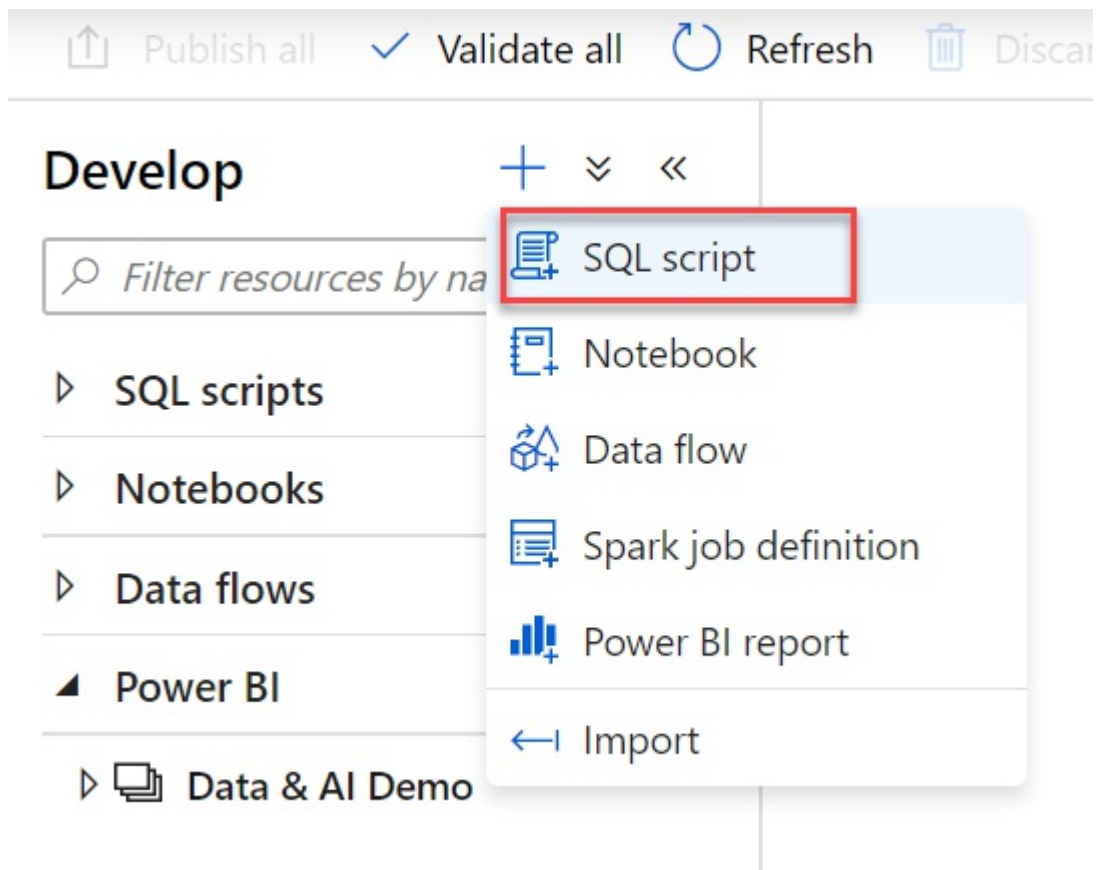
Often in a data warehouse scenario you have users who need their queries to run quickly. The user could be executives of the company who need to run reports or the user could be an analyst running an adhoc query.

Setting importance in Synapse SQL for Azure Synapse allows you to influence the scheduling of queries. Queries with higher importance will be scheduled to run before queries with lower importance. To assign importance to queries, you need to create a workload classifier.

1. Open Synapse Analytics Studio (<https://web.azuresynapse.net/> (<https://web.azuresynapse.net/>)), and then navigate to the **Develop** hub.



2. From the **Develop** menu, select the + button and choose **SQL Script** from the context menu.



3. In the toolbar menu, connect to the **SQL Pool** database to execute the query.



4. In the query window, replace the script with the following to confirm that there are no queries currently being run by users logged in as `asa.sql.workload01`, representing the CEO of the organization or `asa.sql.workload02` representing the data analyst working on the project:

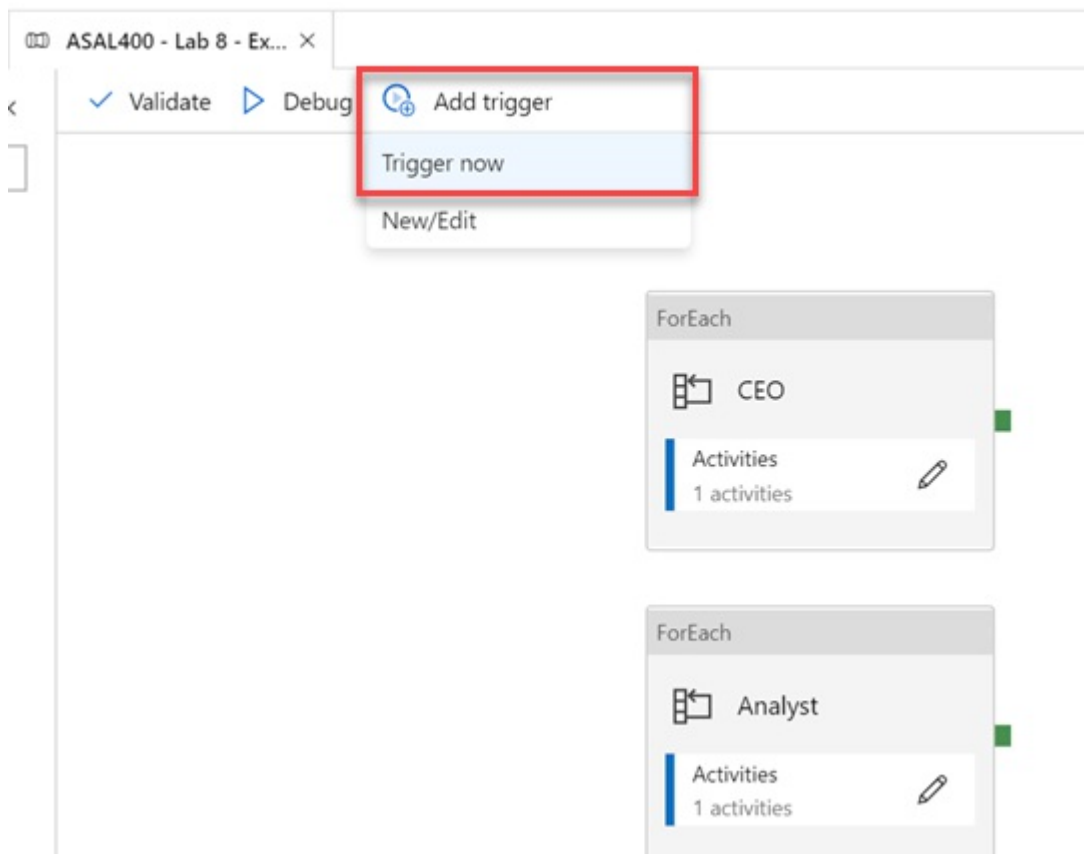
```
--First, let's confirm that there are no queries currently being run by users logged in workload01
or workload02

SELECT s.login_name, r.[Status], r.Importance, submit_time,
start_time ,s.session_id FROM sys.dm_pdw_exec_sessions s
JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id
WHERE s.login_name IN ('asa.sql.workload01','asa.sql.workload02') and Importance
is not NULL AND r.[status] in ('Running','Suspended')
--and submit_time>dateadd(minute,-2,getdate())
ORDER BY submit_time ,s.login_name
```




5. Select **Run** from the toolbar menu to execute the SQL command.



6. You will flood the system with queries and see what happens for `asa.sql.workload01` and `asa.sql.workload02`. To do this, we'll run a Azure Synapse Pipeline which triggers queries. Select the Integrate hub. **Run the Lab 08 - Execute Data Analyst and CEO Queries** Pipeline, which will run / trigger the `asa.sql.workload01` and `asa.sql.workload02` queries. You can run the pipeline with the Debug option if you have an instance of the Integration Runtime running. Otherwise, select **Add trigger**, then **Trigger now**. In the dialog that appears, select **OK**.



7. Navigate to the **Monitor** hub, select **Pipeline runs**, and view the status of the current pipeline run. **Let the Pipeline run for 30 seconds to one minute**, then select **Cancel recursive** for the Lab 08 pipeline.

PIPELINE NAME		RUN START ↑↓	DURATION	TRIGGERED BY	STATUS
<input type="checkbox"/> Lab 08 - Execute D...		5/2/20, 11:16:51 PM	00:03:14	Manual trigger	 In progress
<input type="checkbox"/> Lab 08 - Execute Data Analyst		5/2/20, 11:16:51 PM	00:04:14	Manual trigger	 In progress

8. Let's see what happened to all the queries we just triggered as they flood the system. Return to the query window, replace the script with the following:

```
SELECT s.login_name, r.[Status], r.Importance, submit_time, start_time ,s.session_id FROM
sys.dm_pdw_exec_sessions s
JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id
WHERE s.login_name IN ('asa.sql.workload01','asa.sql.workload02') and Importance
is not NULL AND r.[status] in ('Running','Suspended') and submit_time>dateadd(minute,-2,getdate())
ORDER BY submit_time ,status
```

9. Select **Run** from the toolbar menu to execute the SQL command. You should see an output similar to the following:

```

1 SELECT s.login_name, r.[Status], r.Importance, submit_time, start_time ,s.session_id FROM sys.dm_pdw_exec_sessions s
2 JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id WHERE s.login_name IN ('asa.sql.workload01','asa.sql
3 is not NULL AND r.[status] in ('Running','Suspended') and submit_time>dateadd(minute,-2,getdate())
4 ORDER BY submit_time ,status

```

Results Messages

View Table Chart Export results

Search

LOGIN_NAME	STATUS	IMPORTANCE	SUBMIT_TIME	START_TIME	SESSION_ID
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID819
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID821
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID822
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID823
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID825
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID830
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID832
asa.sql.workload02	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID831
asa.sql.workload01	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID834
asa.sql.workload02	Running	normal	2020-04-29T01:0...	2020-04-29T01:0...	SID835

10. We will give our asa.sql.workload01 user queries priority by implementing the **Workload Importance** feature. Return to the query window, replace the script with the following:

```

IF EXISTS (SELECT * FROM sys.workload_management_workload_classifiers WHERE name = 'CEO')
BEGIN
    DROP WORKLOAD CLASSIFIER CEO;
END
CREATE WORKLOAD CLASSIFIER CEO
WITH (WORKLOAD_GROUP = 'largerc'
,MEMBERNAME = 'asa.sql.workload01',IMPORTANCE = High);

```

11. Select **Run** from the toolbar menu to execute the SQL command.
12. Let's flood the system again with queries and see what happens this time for asa.sql.workload01 and asa.sql.workload02 queries. To do this, we'll run an Azure Synapse Pipeline which triggers queries. **Select** the **Integrate** hub, **run** the **Lab 08 - Execute Data Analyst and CEO Queries** Pipeline, which will run / trigger the asa.sql.workload01 and asa.sql.workload02 queries. **Please let this run for 30 seconds to one minute, then cancel recursively just as you have done before.**
13. In the query window, replace the script with the following to see what happens to the asa.sql.workload01 queries this time:

```

SELECT s.login_name, r.[Status], r.Importance, submit_time, start_time ,s.session_id FROM
sys.dm_pdw_exec_sessions s
JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id
WHERE s.login_name IN ('asa.sql.workload01','asa.sql.workload02') and Importance
is not NULL AND r.[status] in ('Running','Suspended') and submit_time>dateadd(minute,-2,getdate())
ORDER BY submit_time ,status desc

```

14. Select **Run** from the toolbar menu to execute the SQL command. You should see an output similar to the following that shows query executions for the asa.sql.workload01 user having a **high** importance.

```

1 SELECT s.login_name, r.[Status], r.Importance, submit_time, start_time, s.session_id FROM sys.dm_pdw_exec_sessions s
2 JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id WHERE s.login_name IN ('asa.sql.workload01','asa.sql.
3 is not NULL AND r.[status] in ('Running','Suspended') and submit_time>dateadd(minute,-2,getdate())
4 ORDER BY submit_time ,status desc

```

Results

Messages

View

Table

Chart

Export results

Search

LOGIN_NAME	STATUS	IMPORTANCE	SUBMIT_TIME	START_TIME	SESSION_ID
asa.sql.workload01	Running	high	2020-04-29T01:3...	2020-04-29T01:3...	SID907
asa.sql.workload02	Suspended	normal	2020-04-29T01:3...	NULL	SID911
asa.sql.workload02	Suspended	normal	2020-04-29T01:3...	NULL	SID912
asa.sql.workload02	Suspended	normal	2020-04-29T01:3...	NULL	SID913
asa.sql.workload01	Running	high	2020-04-29T01:3...	2020-04-29T01:3...	SID917
asa.sql.workload01	Suspended	high	2020-04-29T01:3...	NULL	SID920
asa.sql.workload02	Suspended	normal	2020-04-29T01:3...	NULL	SID922
asa.sql.workload01	Suspended	high	2020-04-29T01:3...	NULL	SID923
asa.sql.workload01	Suspended	high	2020-04-29T01:3...	NULL	SID924
asa.sql.workload01	Suspended	high	2020-04-29T01:3...	NULL	SID929
asa.sql.workload02	Suspended	normal	2020-04-29T01:3...	NULL	SID930

## Task 2 - Workload Isolation

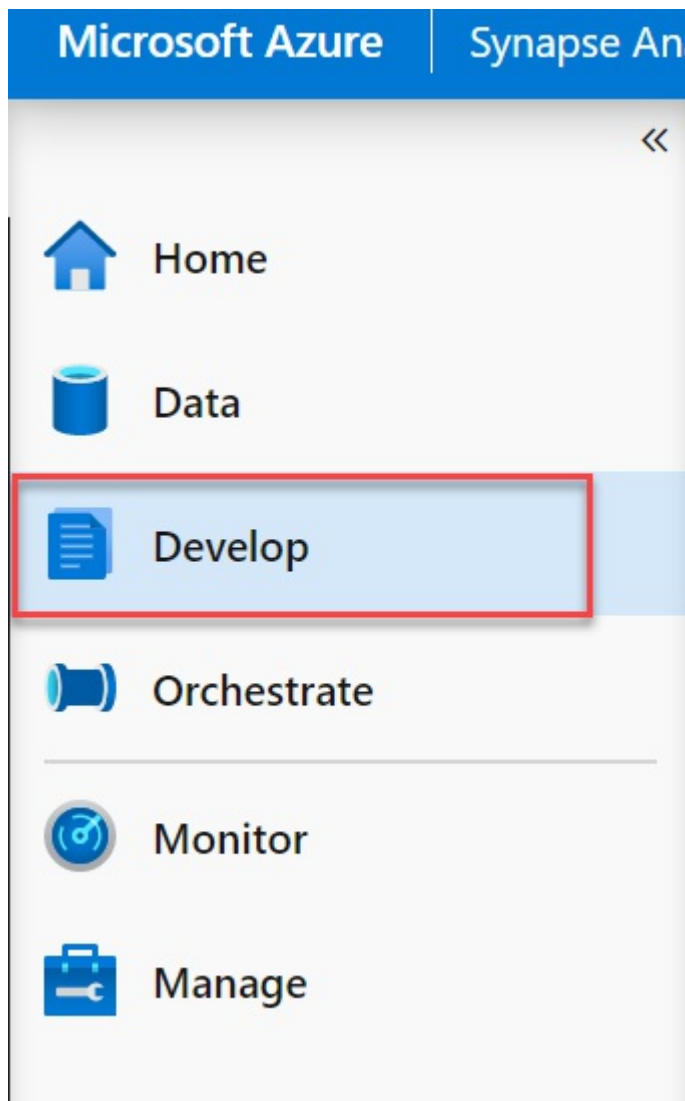
Workload isolation means resources are reserved, exclusively, for a workload group. Workload groups are containers for a set of requests and are the basis for how workload management, including workload isolation, is configured on a system. A simple workload management configuration can manage data loads and user queries.

In the absence of workload isolation, requests operate in the shared pool of resources. Access to resources in the shared pool is not guaranteed and is assigned on an importance basis.


Configuring workload isolation should be done with caution as the resources are allocated to the workload group even if there are no active requests in the workload group. Over-configuring isolation can lead to diminished overall system utilization.

Users should avoid a workload management solution that configures 100% workload isolation: 100% isolation is achieved when the sum of `min_percentage_resource` configured across all workload groups equals 100%. This type of configuration is overly restrictive and rigid, leaving little room for resource requests that are accidentally misclassified. There is a provision to allow one request to execute from workload groups not configured for isolation.

1. Open Synapse Analytics Studio (<https://web.azuresynapse.net/>), and then navigate to the **Develop** hub.



2. From the **Develop** menu, select the + button and choose **SQL Script** from the context menu.



3. In the toolbar menu, connect to the **SQL Pool** database to execute the query.



4. In the query window, replace the script with the following:

```
IF NOT EXISTS (SELECT * FROM sys.workload_management_workload_groups where name = 'CEODemo')
BEGIN
    Create WORKLOAD GROUP CEODemo WITH
    ( MIN_PERCENTAGE_RESOURCE = 50 -- integer value
    ,REQUEST_MIN_RESOURCE_GRANT_PERCENT = 25 --
    ,CAP_PERCENTAGE_RESOURCE = 100
    )
END
```

The code creates a workload group called CEODemo to reserve resources exclusively for the workload group. In this example, a workload group with a MIN\_PERCENTAGE\_RESOURCE set to 50% and REQUEST\_MIN\_RESOURCE\_GRANT\_PERCENT set to 25% is guaranteed 2 concurrency.

5. Select **Run** from the toolbar menu to execute the SQL command.
6. In the query window, replace the script with the following to create a workload Classifier called CEODreamDemo that assigns a workload group and importance to incoming requests:

```
IF NOT EXISTS (SELECT * FROM sys.workload_management_workload_classifiers where name =
'CEODreamDemo')
BEGIN
    Create Workload Classifier CEODreamDemo with
    ( Workload_Group = 'CEODemo',MemberName='asa.sql.workload02',IMPORTANCE = BELOW_NORMAL);
END
```

7. Select **Run** from the toolbar menu to execute the SQL command.
8. In the query window, replace the script with the following to confirm that there are no active queries being run by asa.sql.workload02:



```
SELECT s.login_name, r.[Status], r.Importance, submit_time,
start_time ,s.session_id FROM sys.dm_pdw_exec_sessions s
JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id
WHERE s.login_name IN ('asa.sql.workload02') and Importance
is not NULL AND r.[status] in ('Running','Suspended')
ORDER BY submit_time, status
```

**Note:** There still may be queries left over from previous activities. You can choose to wait for them to complete, or access the SQL requests in the **Manage Hub** and cancel each session manually. Queries that are queued after the creation of the workload group and workload classifier will automatically have this workload management assigned to them.

- Let's flood the system with queries and see what happens for asa.sql.workload02. To do this, we will run an Azure Synapse Pipeline which triggers queries. Select the **Integrate** hub. **Run the Lab 08 - Execute Business Analyst Queries** Pipeline, which will run / trigger asa.sql.workload02 queries. Please let this pipeline run for 30 seconds to one minute, then cancel it recursively.
- In the query window, replace the script with the following to see what happened to all the asa.sql.workload02 queries we just triggered as they flood the system:

```
SELECT s.login_name, r.[Status], r.Importance, submit_time,
start_time ,s.session_id FROM sys.dm_pdw_exec_sessions s
JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id
WHERE s.login_name IN ('asa.sql.workload02') and Importance
is not NULL AND r.[status] in ('Running','Suspended')
ORDER BY submit_time, status
```

- Select **Run** from the toolbar menu to execute the SQL command. You should see an output similar to the following that shows the importance for each session set to below\_normal:

```
1 SELECT s.login_name, r.[Status], r.Importance, submit_time,
2 start_time ,s.session_id FROM sys.dm_pdw_exec_sessions s
3 JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id
4 WHERE s.login_name IN ('asa.sql.workload02') and Importance
5 is not NULL AND r.[status] in ('Running','Suspended')
6 ORDER BY submit_time, status
```

Results Messages					
View Table Chart Export results					
Search					
LOGIN_NAME	STATUS	IMPORTANCE	SUBMIT_TIME	START_TIME	SESSION_ID
asa.sql.workload02	Running	below_normal	2020-04-29T02:3...	2020-04-29T02:3...	SID996
asa.sql.workload02	Running	below_normal	2020-04-29T02:3...	2020-04-29T02:3...	SID999
asa.sql.workload02	Suspended	below_normal	2020-04-29T02:3...	NULL	SID1000
asa.sql.workload02	Suspended	below_normal	2020-04-29T02:3...	NULL	SID1007
asa.sql.workload02	Suspended	below_normal	2020-04-29T02:3...	NULL	SID1006
asa.sql.workload02	Suspended	below_normal	2020-04-29T02:3...	NULL	SID1008
asa.sql.workload02	Suspended	below_normal	2020-04-29T02:3...	NULL	SID1009
asa.sql.workload02	Suspended	below_normal	2020-04-29T02:3...	NULL	SID1012

- In the query window, replace the script with the following to set 3.25% minimum resources per request:

```

IF EXISTS (SELECT * FROM sys.workload_management_workload_classifiers where group_name =
'CEODemo')
BEGIN
    Drop Workload Classifier CEODreamDemo
    DROP WORKLOAD GROUP CEODemo
    --- Creates a workload group 'CEODemo'.
    Create WORKLOAD GROUP CEODemo WITH
    (MIN_PERCENTAGE_RESOURCE = 26 -- integer value
    ,REQUEST_MIN_RESOURCE_GRANT_PERCENT = 3.25 -- factor of 26 (guaranteed more than 4
concurrencies)
    ,CAP_PERCENTAGE_RESOURCE = 100
    )
    --- Creates a workload Classifier 'CEODreamDemo'.
    Create Workload Classifier CEODreamDemo with
    (Workload_Group = 'CEODemo', MemberName='asa.sql.workload02', IMPORTANCE = BELOW_NORMAL);
END

```

**Note:** Configuring workload containment implicitly defines a maximum level of concurrency. With a CAP\_PERCENTAGE\_RESOURCE set to 60% and a REQUEST\_MIN\_RESOURCE\_GRANT\_PERCENT set to 1%, up to a 60-concurrency level is allowed for the workload group. Consider the method included below for determining the maximum concurrency:

$$[\text{Max Concurrency}] = [\text{CAP\_PERCENTAGE\_RESOURCE}] / [\text{REQUEST\_MIN\_RESOURCE\_GRANT\_PERCENT}]$$

- Let's flood the system again and see what happens for asa.sql.workload02. To do this, we will run an Azure Synapse Pipeline which triggers queries. Select the Integrate hub. **Run the Lab 08 - Execute Business Analyst Queries Pipeline**, which will run / trigger asa.sql.workload02 queries. **Please let this run for 30 seconds to one minute, then cancel it recursively.**
- In the query window, replace the script with the following to see what happened to all of the asa.sql.workload02 queries we just triggered as they flood the system:

```

SELECT s.login_name, r.[Status], r.Importance, submit_time,
start_time ,s.session_id FROM sys.dm_pdw_exec_sessions s
JOIN sys.dm_pdw_exec_requests r ON s.session_id = r.session_id
WHERE s.login_name IN ('asa.sql.workload02') and Importance
is not NULL AND r.[status] in ('Running', 'Suspended')
ORDER BY submit_time, status

```

- Select **Run** from the toolbar menu to execute the SQL command.

## Exercise 2 - Workload Monitoring


Azure Synapse Analytics provides a rich monitoring experience within the Azure portal to surface insights regarding your data warehouse workload. The Azure portal is the recommended tool when monitoring your data warehouse as it provides configurable retention periods, alerts, recommendations, and customizable charts and dashboards for metrics and logs. The portal also enables you to integrate with other Azure monitoring services such as Azure Monitor (logs) with Log analytics to provide a holistic monitoring experience for not only your data warehouse but also your entire Azure analytics platform for an integrated monitoring experience.

For a programmatic experience when monitoring SQL Analytics via T-SQL, the service provides a set of Dynamic Management Views (DMVs). These views are useful when actively troubleshooting and identifying performance bottlenecks with your workload.


### Task 1 - Monitoring with Dynamic Management Views

All logins to your data warehouse are logged to sys.dm\_pdw\_exec\_sessions. This DMV contains the last 10,000 logins. The session\_id is the primary key and is assigned sequentially for each new login.

- Open Synapse Analytics Studio (<https://web.azuresynapse.net/> (<https://web.azuresynapse.net/>)), and then navigate to the **Develop** hub.



2. From the **Develop** menu, select the + button and choose **SQL Script** from the context menu.



3. In the toolbar menu, connect to the **SQL Pool** database to execute the query.



4. In the query window, replace the script with the following:

```
SELECT * FROM sys.dm_pdw_exec_sessions where status <> 'Closed' and session_id <> session_id();
```

All queries executed on SQL pool are logged to sys.dm\_pdw\_exec\_requests. This DMV contains the last 10,000 queries executed. The request\_id uniquely identifies each query and is the primary key for this DMV. The request\_id is assigned sequentially for each new query and is prefixed with QID, which stands for query ID. Querying this DMV for a given session\_id shows all queries for a given login.

5. Select **Run** from the toolbar menu to execute the SQL command.
6. Let's flood the system with queries to create operations to monitor. To do this, we will run a Azure Synapse Pipeline which triggers queries. Select the *Integrate* hub. **Run the Lab 08 - Execute Business Analyst Queries** Pipeline, which will run / trigger asa.sql.workload02 queries. **Please let this pipeline run for 30 seconds to one minute, then cancel it recursively.**
7. In the query window, replace the script with the following:

```
SELECT *
FROM sys.dm_pdw_exec_requests
WHERE status not in ('Completed', 'Failed', 'Cancelled')
AND session_id <> session_id()
ORDER BY submit_time DESC;
```

8. Select **Run** from the toolbar menu to execute the SQL command. You should see a list of sessions in the query results similar to the following. **Note the Request\_ID of a query** in the results that you would like to investigate (*keep this value in a text editor for a later step*):

```

1 SELECT *
2 FROM sys.dm_pdw_exec_requests
3 WHERE status not in ('Completed', 'Failed', 'Cancelled')
4 AND session_id <> session_id()
5 ORDER BY submit_time DESC;

```

Results Messages

View Table Chart Export results

Search

REQUEST_ID	SESSION_ID	STATUS	SUBMIT_TIME	START_TIME	END_COMPILE_TIME	END_TIME	TOTAL_ELAPSED_TII
QID6386	SID1070	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5375
QID6388	SID1071	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5375
QID6390	SID1072	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5375
QID6379	SID1065	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5484
QID6381	SID1066	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5484
QID6373	SID1062	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5500
QID6376	SID1064	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5500
QID6377	SID1063	Running	2020-04-29T02:5...	2020-04-29T02:5...	2020-04-29T02:5...	NULL	5500

9. As an alternative, you can execute the following SQL command to find the top 10 longest running queries.

```

SELECT TOP 10 *
FROM sys.dm_pdw_exec_requests
ORDER BY total_elapsed_time DESC;

```

10. To simplify the lookup of a query in the `sys.dm_pdw_exec_requests` table, use `LABEL` to assign a comment to your query, which can be looked up in the `sys.dm_pdw_exec_requests` view. To test using the labels, replace the script in the query window with the following:

```

SELECT *
FROM sys.tables
OPTION (LABEL = 'My Query');

```

11. Select **Run** from the toolbar menu to execute the SQL command.

12. In the query window, replace the script with the following to filter the results with the label, `My Query`.

```

-- Find a query with the Label 'My Query'
-- Use brackets when querying the label column, as it is a key word
SELECT *
FROM sys.dm_pdw_exec_requests
WHERE [label] = 'My Query';

```

13. Select **Run** from the toolbar menu to execute the SQL command. You should see the previously run query in the results view.

14. In the query window, replace the script with the following to retrieve the query's distributed SQL (DSQL) plan from `sys.dm_pdw_request_steps`. **Be sure to replace** the `QID####` with the `Request_ID` you noted in Step 8:

```

SELECT * FROM sys.dm_pdw_request_steps
WHERE request_id = 'QID####'
ORDER BY step_index;

```

15. Select **Run** from the toolbar menu to execute the SQL command. You should see results showing the distributed query plan steps for the specified request:

```

1 SELECT * FROM sys.dm_pdw_request_steps
2 WHERE request_id = 'QID6377'
3 ORDER BY step_index;

```

Results Messages

View Table Chart Export results

Search

REQUEST_ID	STEP_INDEX	OPERATION_TYPE	DISTRIBUTION_TYPE	LOCATION_TYPE	STATUS	ERROR_ID
QID6377	0	OnOperation	Unspecified	Control	Complete	NULL
QID6377	1	PartitionMoveOperation	AllDistributions	Compute	Complete	NULL
QID6377	2	ReturnOperation	Unspecified	Control	Complete	NULL
QID6377	3	OnOperation	Unspecified	Control	Complete	NULL

When a DSQL plan is taking longer than expected, the cause can be a complex plan with many DSQL steps or just one step taking a long time. If the plan is many steps with several move operations, consider optimizing your table distributions to reduce data movement.

## Task 2 - Orchestration Monitoring with the Monitor Hub

- Let's run a pipeline to monitor its execution in the next step. To do this, select the **Integrate** hub. Run the **Lab 08 - Execute Business Analyst Queries** Pipeline.

ASAL400 - Lab 8 - Ex... ASAL400 - Lab 8 - Ex...

Activities

Search activities

- Synapse
- Move & transform
- Azure Data Explorer
- Azure Function
- Batch Service
- Databricks
- Data Lake Analytics
- General
- HDInsight
- Iteration & conditionals
- Machine Learning

Validate Debug Add trigger

Trigger now

New/Edit

2. Navigate to the Monitor hub. Then select **Pipeline runs** to get a list of pipelines that ran during the last 24 hours. Observe the Pipeline status.

The screenshot shows the Microsoft Azure Synapse Analytics interface. On the left, the 'Monitor' hub is selected. In the 'Orchestration' section, 'Pipeline runs' is highlighted. The main panel displays a table of pipeline runs. The first row shows a pipeline named 'Lab 08 - Execute Business Analyst Queries' with a status of 'In progress'.

PIPELINE NAME	RUN START	DURATION	TRIGGERED BY	STATUS	PARAMETERS
Lab 08 - Execute Business Analyst Queries	8/14/20, 9:58:32 PM	00:04:05	Manual trigger	In progress	

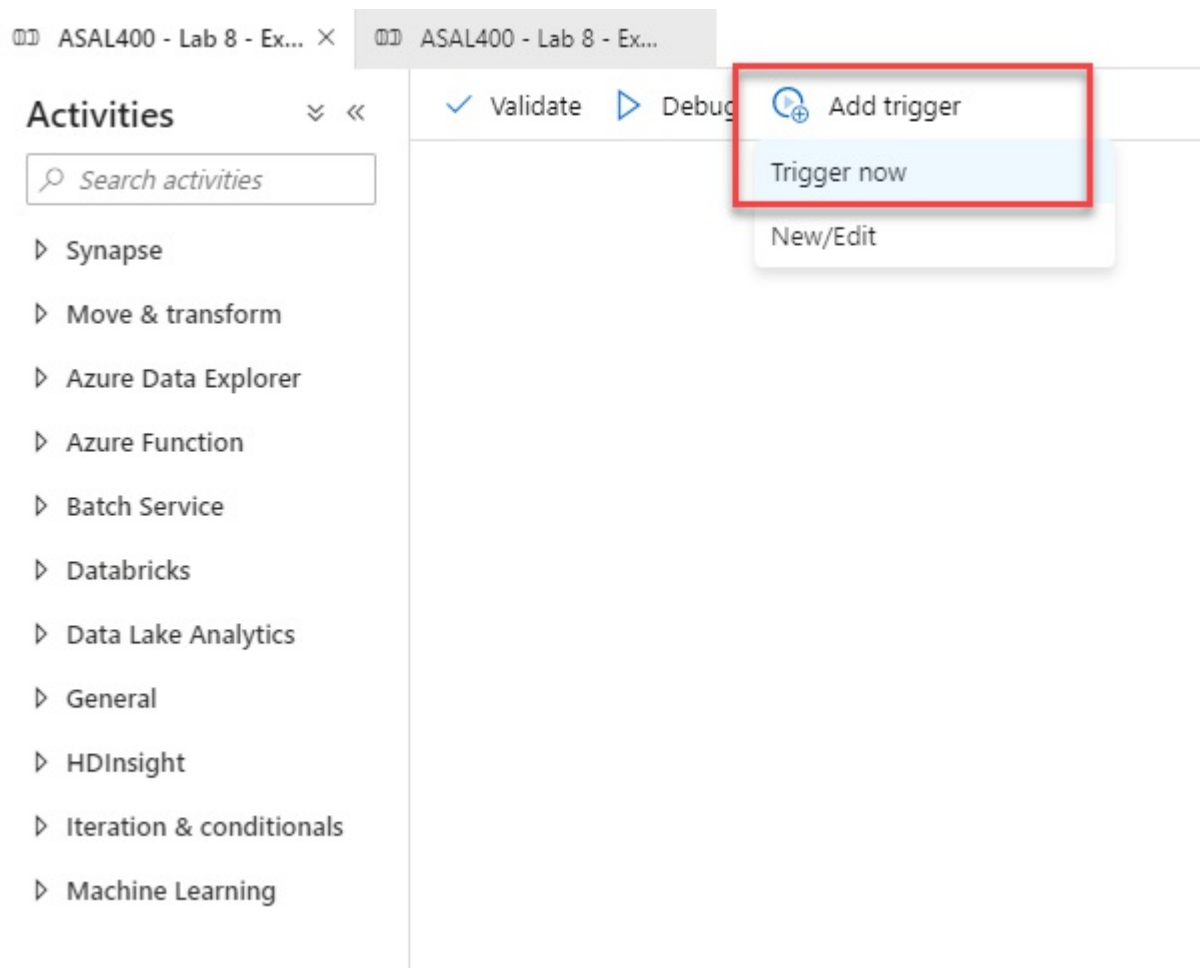
3. Hover over the running pipeline and select **Cancel** to cancel the execution of the current instance of the pipeline.

This screenshot is a closer view of the 'Pipeline runs' table. The first row is highlighted, and the 'Cancel' button (a circle with a diagonal line) is visible next to the pipeline name. A tooltip labeled 'Cancel' appears when the button is hovered over.


PIPELINE NAME	RUN START	DURATION	TRIGGERED BY
Lab 08 - Execute Business Analyst Queries	8/14/20, 9:58:32 PM	00:04:05	Manual trigger
Lab 08 - Execute Business Analyst Queries	8/14/20, 9:49:37 PM	00:05:16	Manual trigger
Lab 08 - Execute Business Analyst Queries	8/14/20, 9:35:50 PM	00:01:13	Manual trigger

### Task 3 - Monitoring SQL Requests with the Monitor Hub

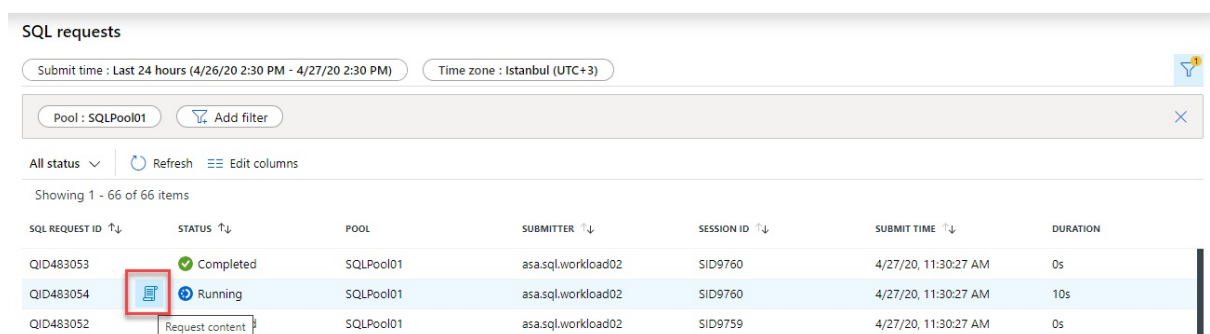
1. Let's run a pipeline to monitor its execution in the next step. To do this, select the **Integrate** hub. Run the **Lab 08 - Execute Business Analyst Queries** Pipeline. Let this pipeline run for 30 seconds to one minute, then cancel it recursively.



2. Navigate to the Monitor hub. Then select **SQL requests** to get a list of SQL requests that ran during the last 24 hours.
3. Select the **Pool** filter and select your SQL Pool. Observe the Request Submitter, Submit Time, Duration, and Queued Duration values.



4. Hover onto a SQL Request log and select Request Content to access the actual T-SQL command executed as part of the SQL Request.



## Resources

- [Workload Group Isolation \(https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-isolation\)](https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-isolation)
- [Workload Isolation \(https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-isolation\)](https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-isolation)



- [Workload Importance \(https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-importance\)](https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-importance)
- [Workload Classification \(https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-classification\)](https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-classification)
- [Monitoring workload using DMVs \(https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-manage-monitor\)](https://docs.microsoft.com/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-manage-monitor)